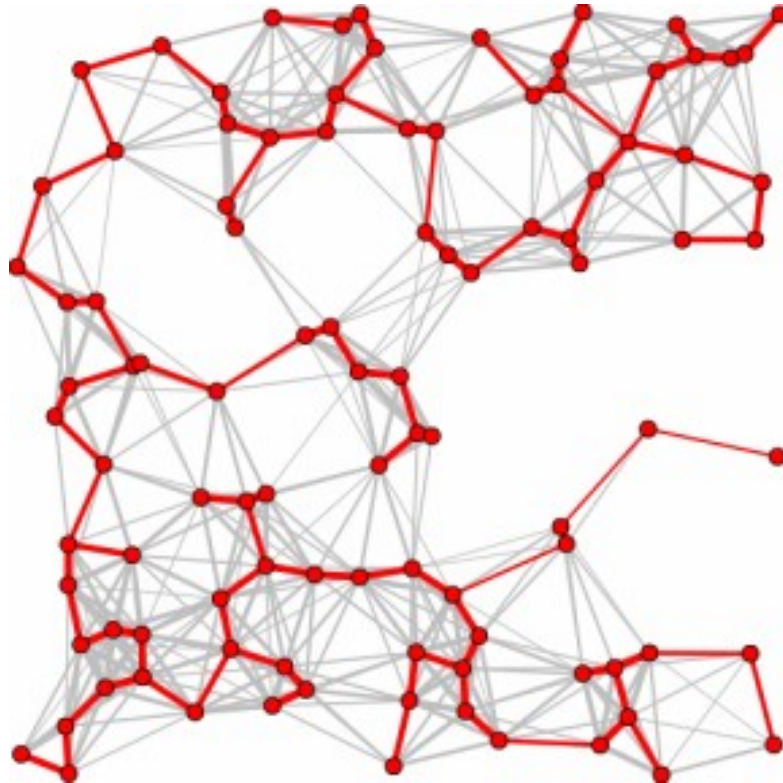


Minimum Spanning Trees



Outline and Reading

- Minimum Spanning Trees (7.3)
 - Definitions
 - A crucial fact
- The Prim-Jarnik Algorithm (7.3.2)
- Kruskal's Algorithm (7.3.1)
- Baruvka's Algorithm (7.3.3)

Minimum Spanning Tree

Spanning subgraph

- Subgraph of a graph G containing all the vertices of G

Spanning tree


- Spanning subgraph that is itself a (free) tree

Minimum spanning tree (MST)

- Spanning tree of a weighted graph with **minimum total edge weight**

Applications

- Communications networks
- Transportation networks




Cycle Property

Cycle Property:



- Let T be a minimum spanning tree of a weighted graph G
- Let e be an edge of G that is not in T and C let be the cycle formed by e with T
- For every edge f of C , $\text{weight}(f) \leq \text{weight}(e)$

Proof:

- By contradiction
- If $\text{weight}(f) > \text{weight}(e)$ we can get a spanning tree of smaller weight by replacing e with f



Replacing f with e
yields a better
spanning tree




Partition Property

Partition Property:


- Consider a partition of the vertices of G into subsets U and V
- Let e be an edge of **minimum weight across the partition**
- There is a minimum spanning tree of G **containing** edge e

Proof:

- Let T be an MST of G
- If T does not contain e , consider the cycle C formed by e with T and let f be an edge of C across the partition
- By the cycle property, $\text{weight}(f) \leq \text{weight}(e)$
- Thus, $\text{weight}(f) = \text{weight}(e)$
- We obtain another MST by replacing f with e



Replacing f with e yields another MST




Prim-Jarnik's Algorithm

Idea:

- Similar to Dijkstra's algorithm (for a connected graph)
- Pick an arbitrary vertex s and we **grow the MST as a cloud** of vertices, starting from s
- Store with each vertex v a label $d(v)$ = the smallest weight of an edge connecting v to a vertex in the cloud

At each step:

- Add to the cloud the vertex u outside the cloud with the smallest label
- Update the labels of the vertices adjacent to u



Prim-Jarnik's Algorithm (cont.)

A **priority queue** stores the vertices outside the cloud

- Key: distance
- Element: vertex

Locator-based methods

- ***insert(k,e)*** returns a locator
- ***replaceKey(l,k)*** changes the key of an item

We store three labels with each vertex:

- Distance
- Parent edge in MST
- Locator in priority queue

Algorithm *PrimJarnikMST(G)*

Q \leftarrow new heap-based priority queue

s \leftarrow a vertex of ***G***

for all ***v*** \in ***G.vertices()***

if ***v = s***

setDistance(v, 0)

else

setDistance(v, ∞)

setParent(v, \emptyset)

l \leftarrow ***Q.insert(getDistance(v), v)***

setLocator(v,l)

while \neg ***Q.isEmpty()***

u \leftarrow ***Q.removeMin()***

for all ***e*** \in ***G.incidentEdges(u)***

z \leftarrow ***G.opposite(u,e)***

r \leftarrow ***weight(e)***


if ***r < getDistance(z)***

setDistance(z,r)


setParent(z,e)

Q.replaceKey(getLocator(z),r)

Example



Example (contd.)



Analysis

- Graph operations
 - Method `incidentEdges` is called once for each vertex
- Label operations
 - We set/get the distance, parent and locator labels of vertex z $O(\deg(z))$ times
 - Setting/getting a label takes $O(1)$ time
- Priority queue operations
 - Each vertex is inserted once into and removed once from the priority queue, where each insertion or removal takes $O(\log n)$ time
 - The key of a vertex w in the priority queue is modified at most $\deg(w)$ times, where each key change takes $O(\log n)$ time
- Prim-Jarnik's algorithm runs in $O((n + m) \log n)$ time provided the graph is represented by the adjacency list structure
 - Recall that $\sum_v \deg(v) = 2m$
- The running time is $O(m \log n)$ since the graph is connected

Kruskal's Algorithm

A priority queue stores the edges outside the cloud

- Key: weight
- Element: edge

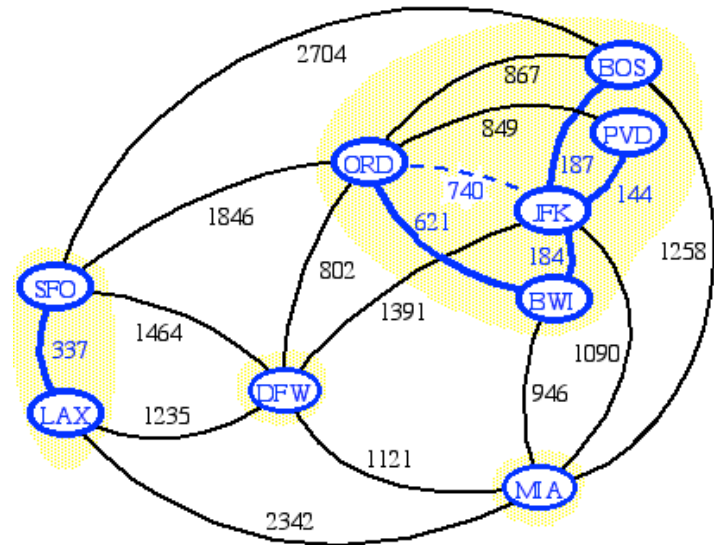
At the end of the algorithm

- We are left with one cloud that encompasses the MST
- A tree T which is our MST


```
Algorithm KruskalMST( $G$ )  
  for each vertex  $V$  in  $G$  do  
    define a Cloud( $v$ ) of  $\leftarrow \{v\}$   
  let  $Q$  be a priority queue  
  Insert all edges into  $Q$  using their weights as the key  
   $T \leftarrow \emptyset$   
  while  $T$  has fewer than  $n-1$  edges do  
    edge  $e = T.removeMin()$   
    Let  $u, v$  be the endpoints of  $e$   
    { check if edge is necessary to connect two clouds }  
    if Cloud( $v$ )  $\neq$  Cloud( $u$ ) then  
      Add edge  $e$  to  $T$   
      Merge Cloud( $v$ ) and Cloud( $u$ )  
  return  $T$ 
```

Data Structure for Kruskal Algorithm

- The algorithm maintains a forest of trees
- An edge is accepted if it connects distinct trees
- We need a data structure that maintains a **partition**, i.e., a collection of disjoint sets, with the operations:
 - **find**(u): return the set storing u
 - **union**(u, v): replace the sets storing u and v with their union



Representation of a Partition



- Each set is stored in a sequence
- Each element has a reference back to the set
 - operation **find**(u) takes $O(1)$ time, and returns the set of which u is a member.
 - in operation **union**(u, v), we move the elements of the smaller set to the sequence of the larger set and update their references
 - the time for operation **union**(u, v) is $\min(n_u, n_v)$, where n_u and n_v are the sizes of the sets storing u and v
- Whenever an element is processed, it goes into a set of size at least double, hence each element is processed at most $\log n$ times

Partition-Based Implementation

A partition-based version of Kruskal's Algorithm performs cloud merges as unions and tests as finds.

Algorithm Kruskal(G):

Input: A weighted graph G .

Output: An MST T for G .

Let P be a partition of the vertices of G , where each vertex forms a separate set.

Let Q be a priority queue storing the edges of G , sorted by their weights

Let T be an initially-empty tree

while Q is not empty **do**

$(u,v) \leftarrow Q.\text{removeMinElement}()$

if $P.\text{find}(u) \neq P.\text{find}(v)$ **then**

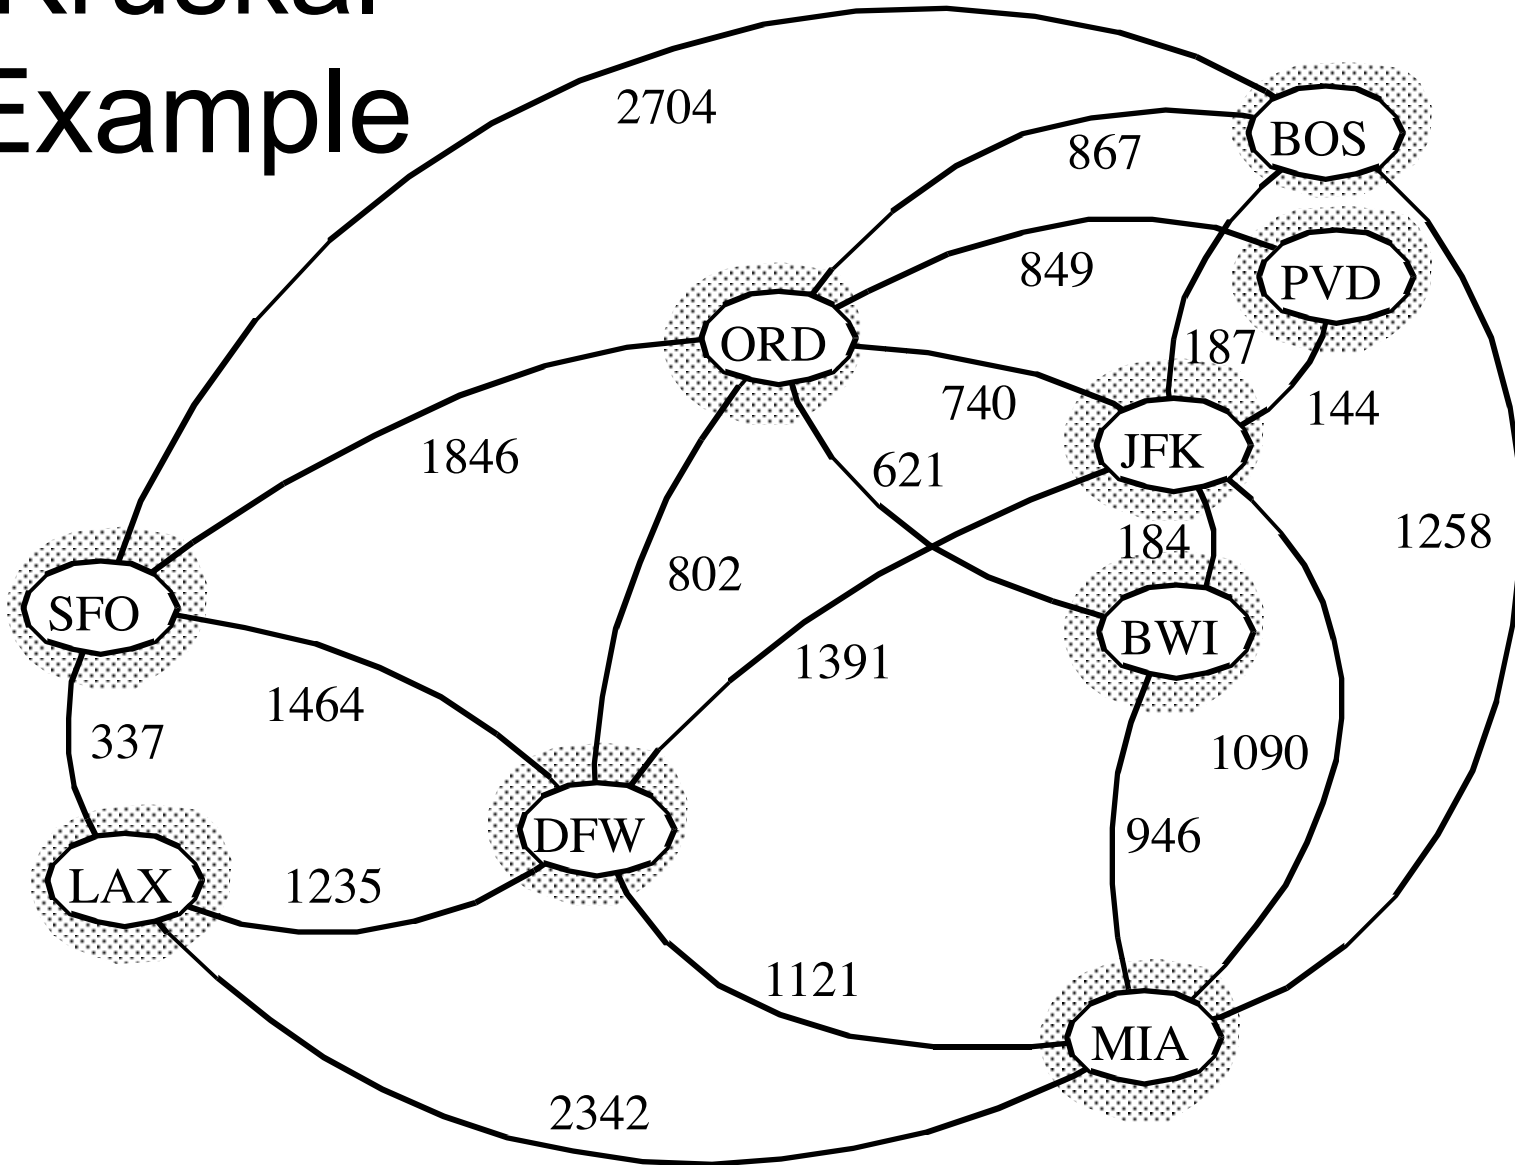
 Add (u,v) to T

$P.\text{union}(u,v)$

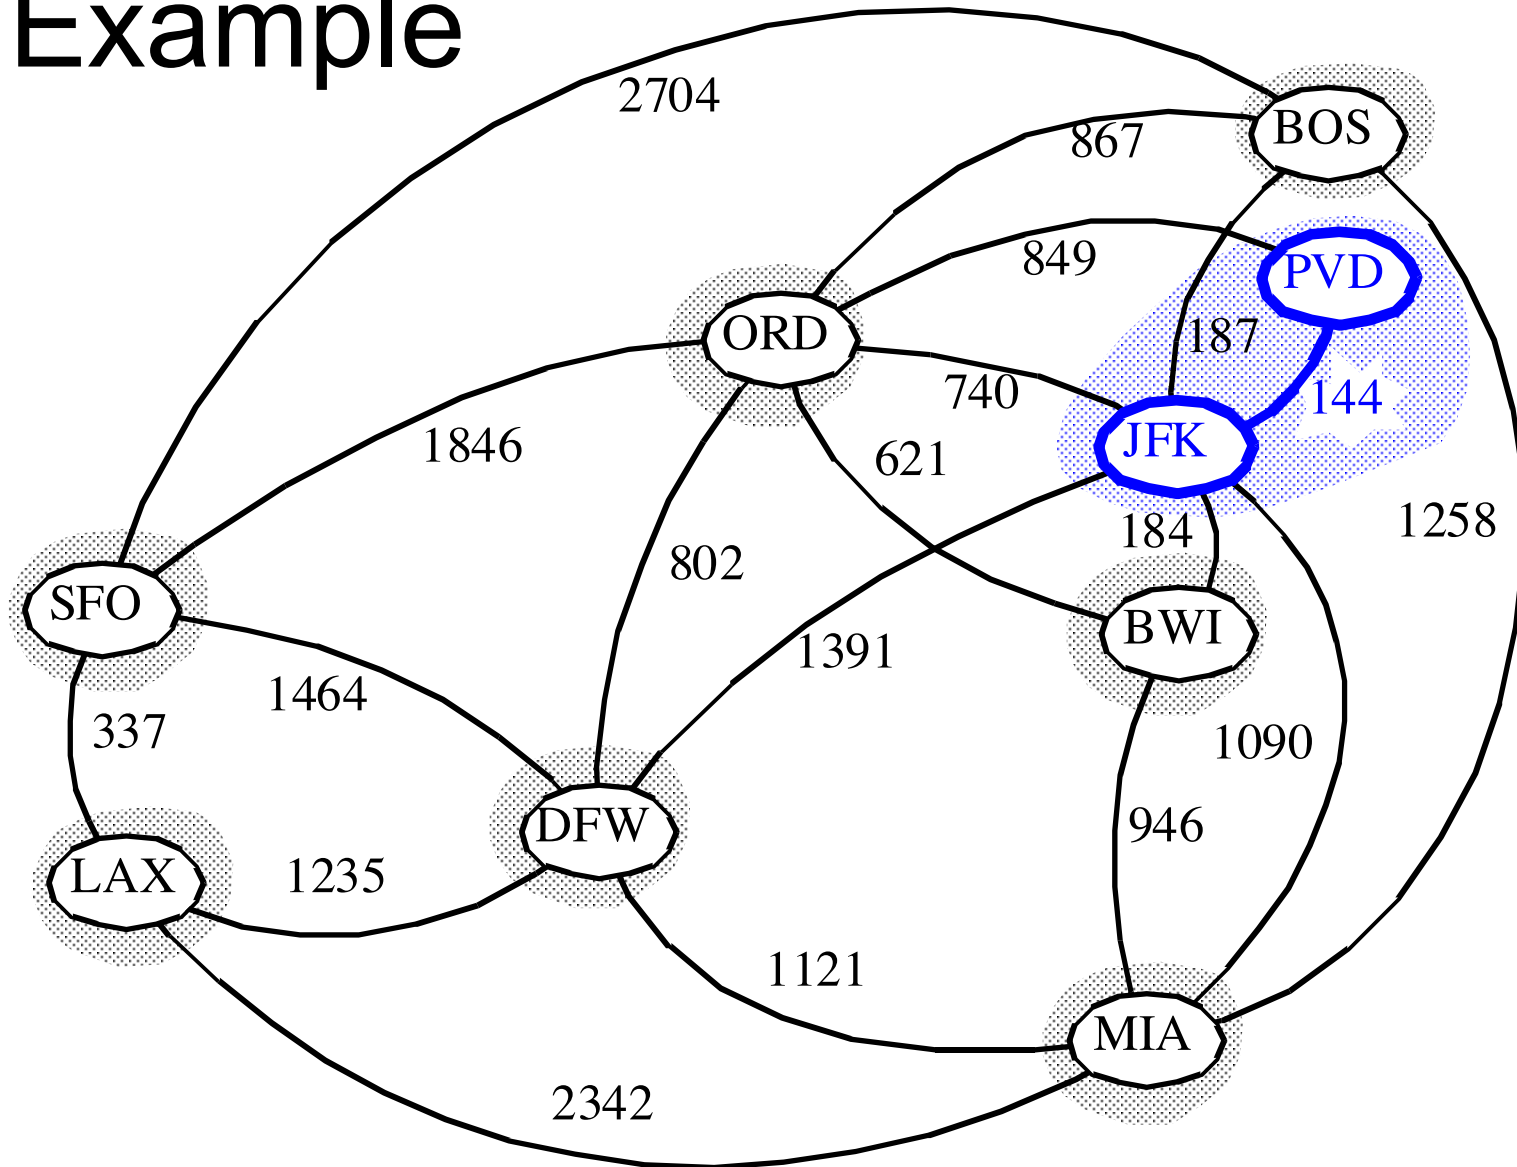
return T

Running time: $O((n+m)\log n)$

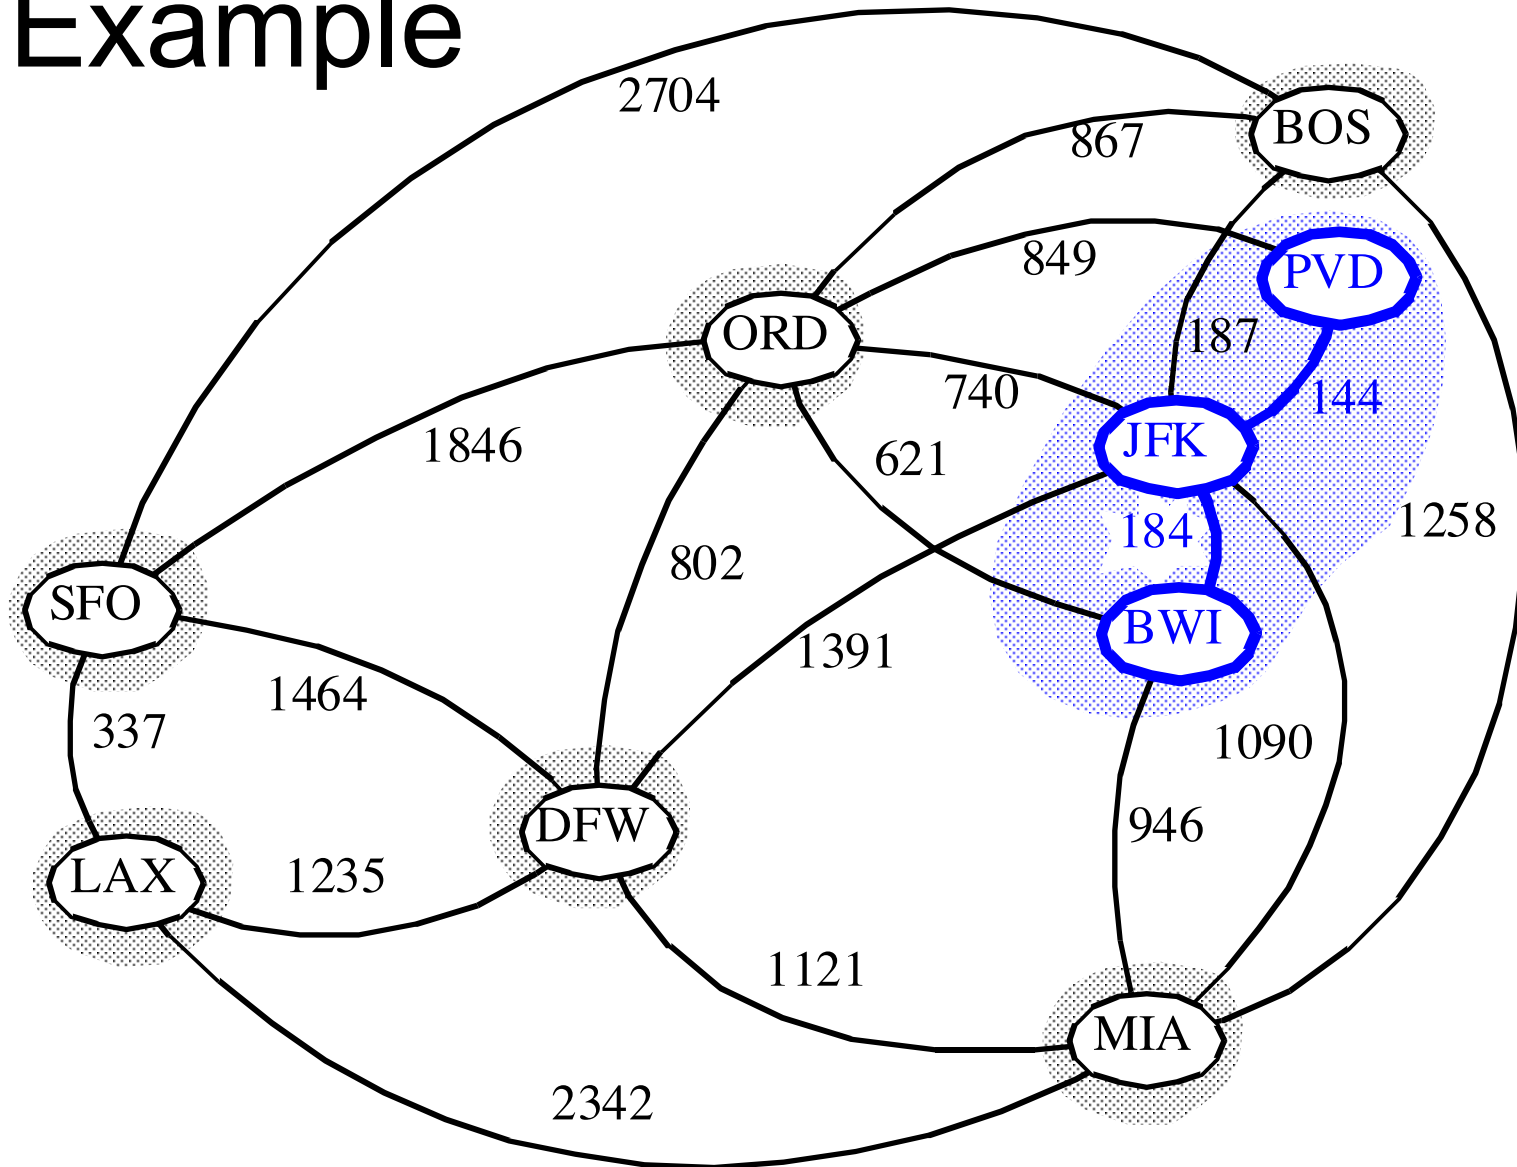
Kruskal Example



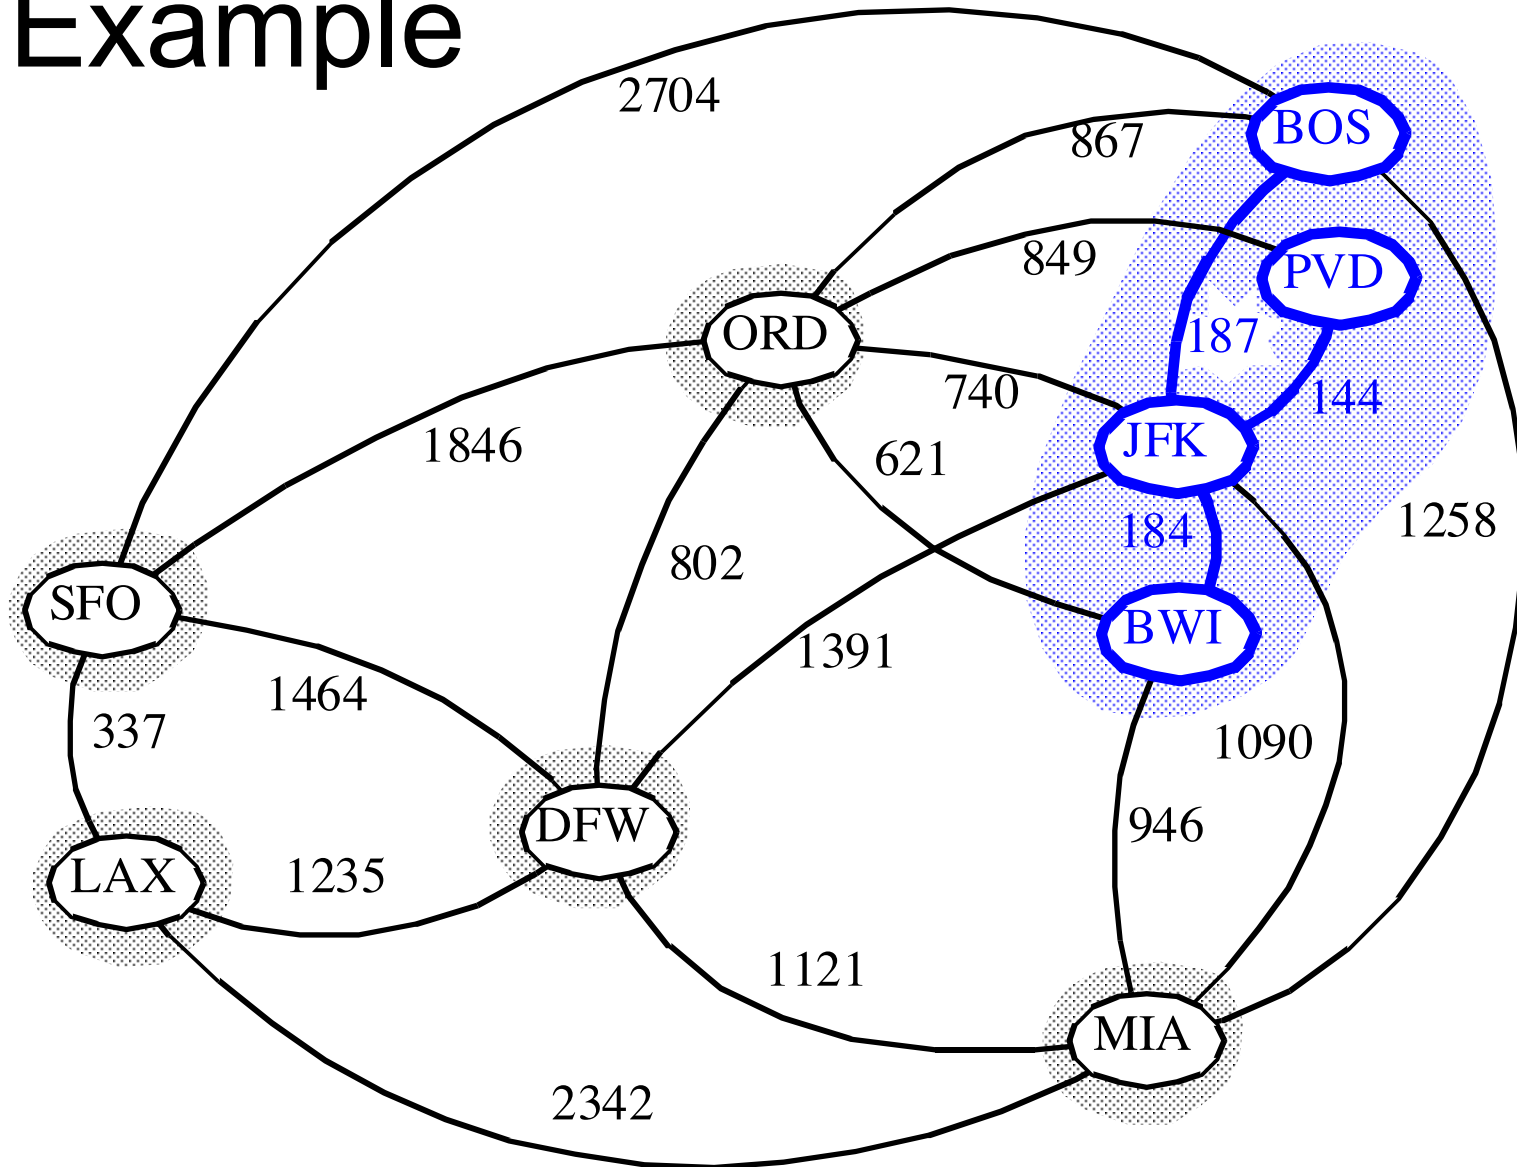
Example



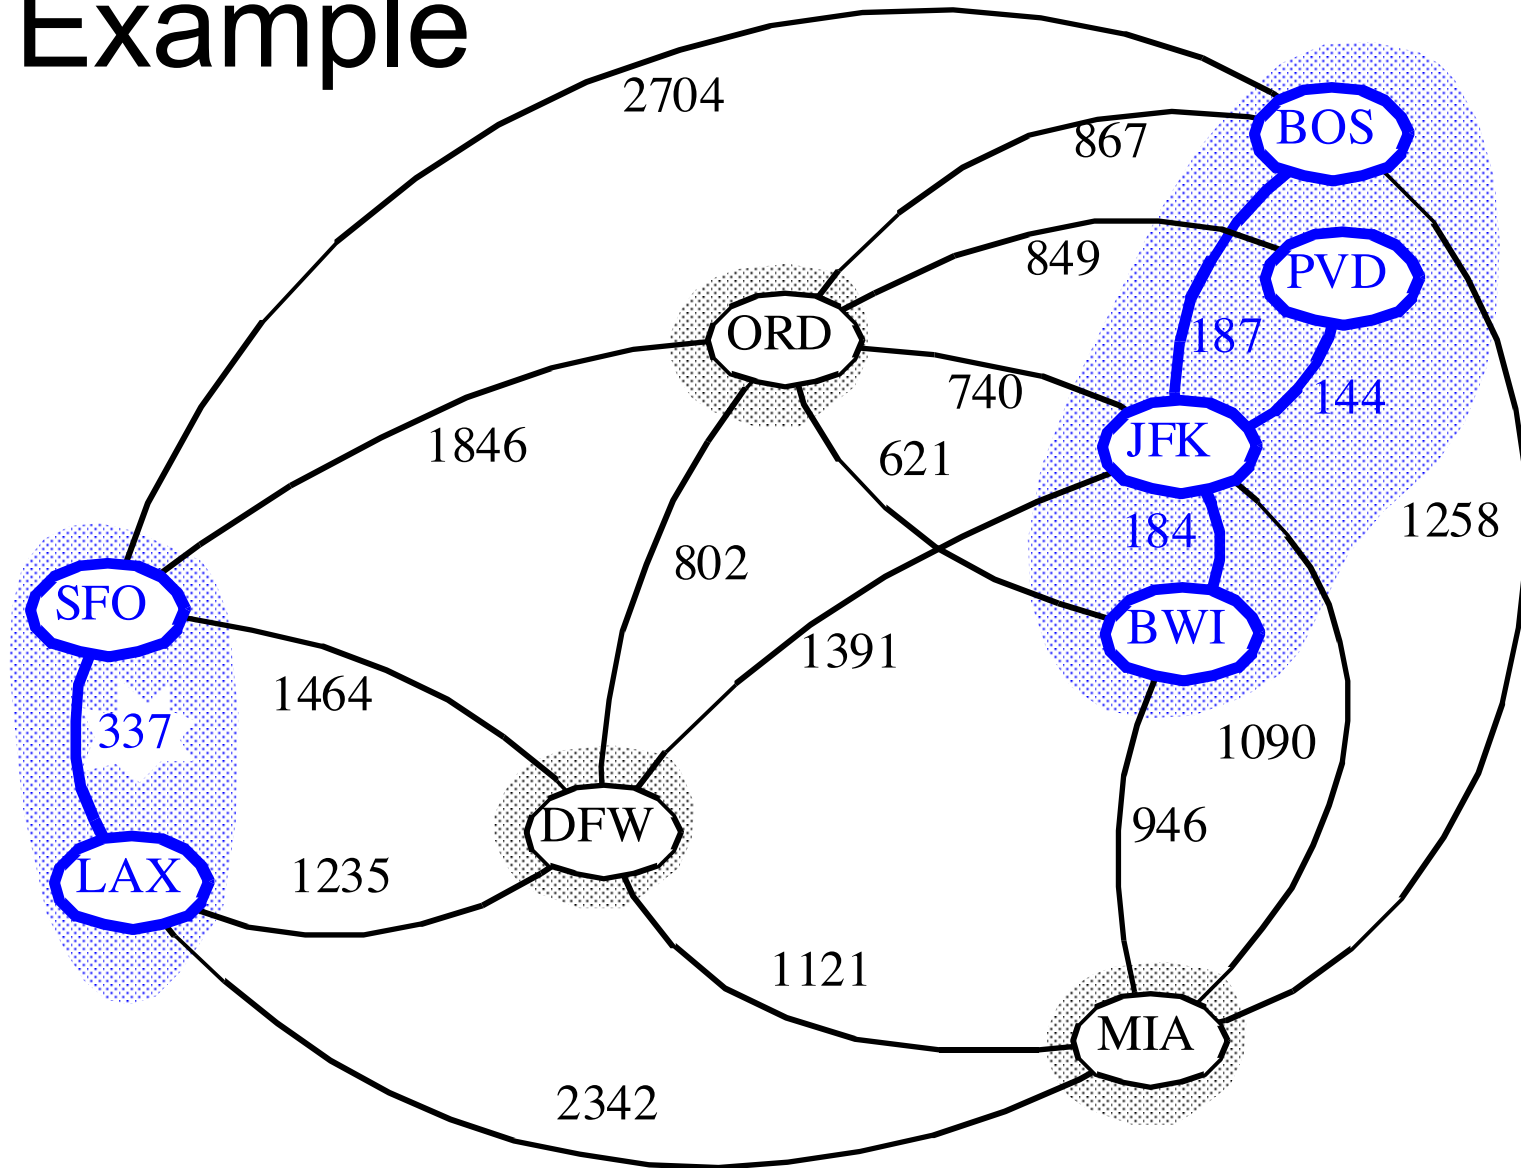
Example



Example



Example



Example

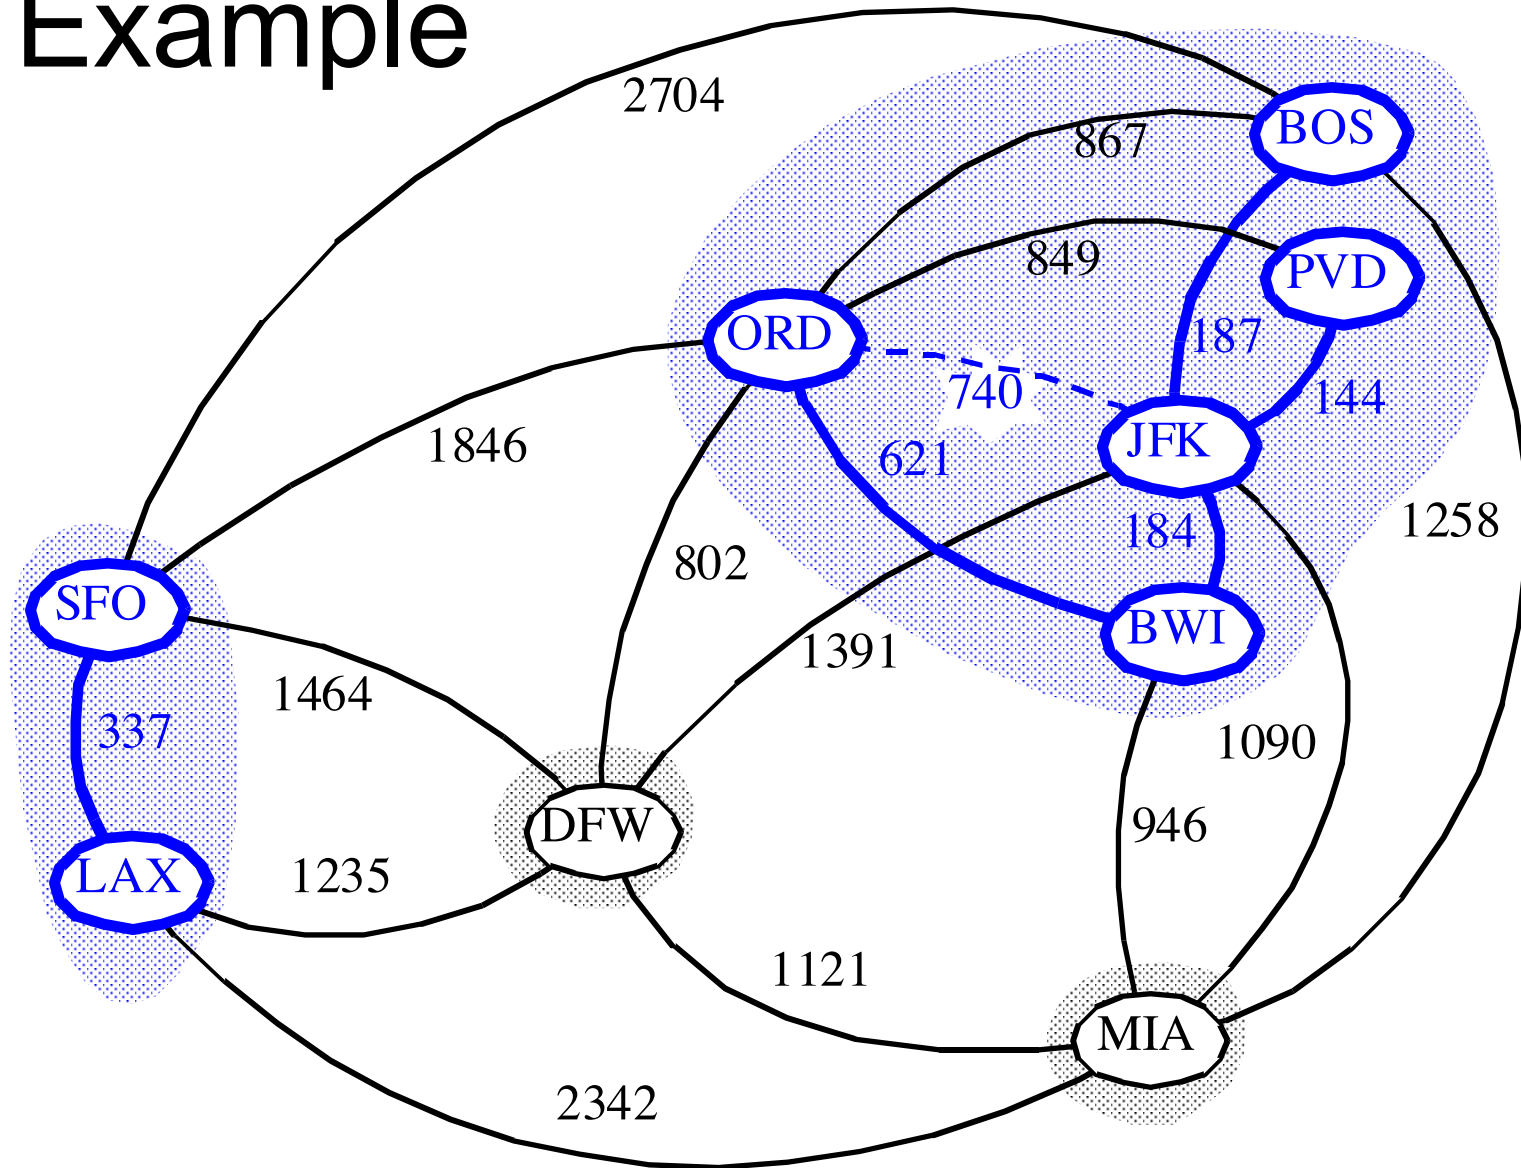
The graph illustrates a network of airports and flight routes. The nodes represent airports: SFO, LAX, DFW, ORD, BOS, PVD, JFK, BWI, and MIA. The edges represent flight routes with associated weights. The highlighted paths are:

- SFO to LAX (weight 337)
- ORD to JFK to BWI (weights 621, 184, 144)
- BOS to PVD to JFK (weights 867, 187, 144)

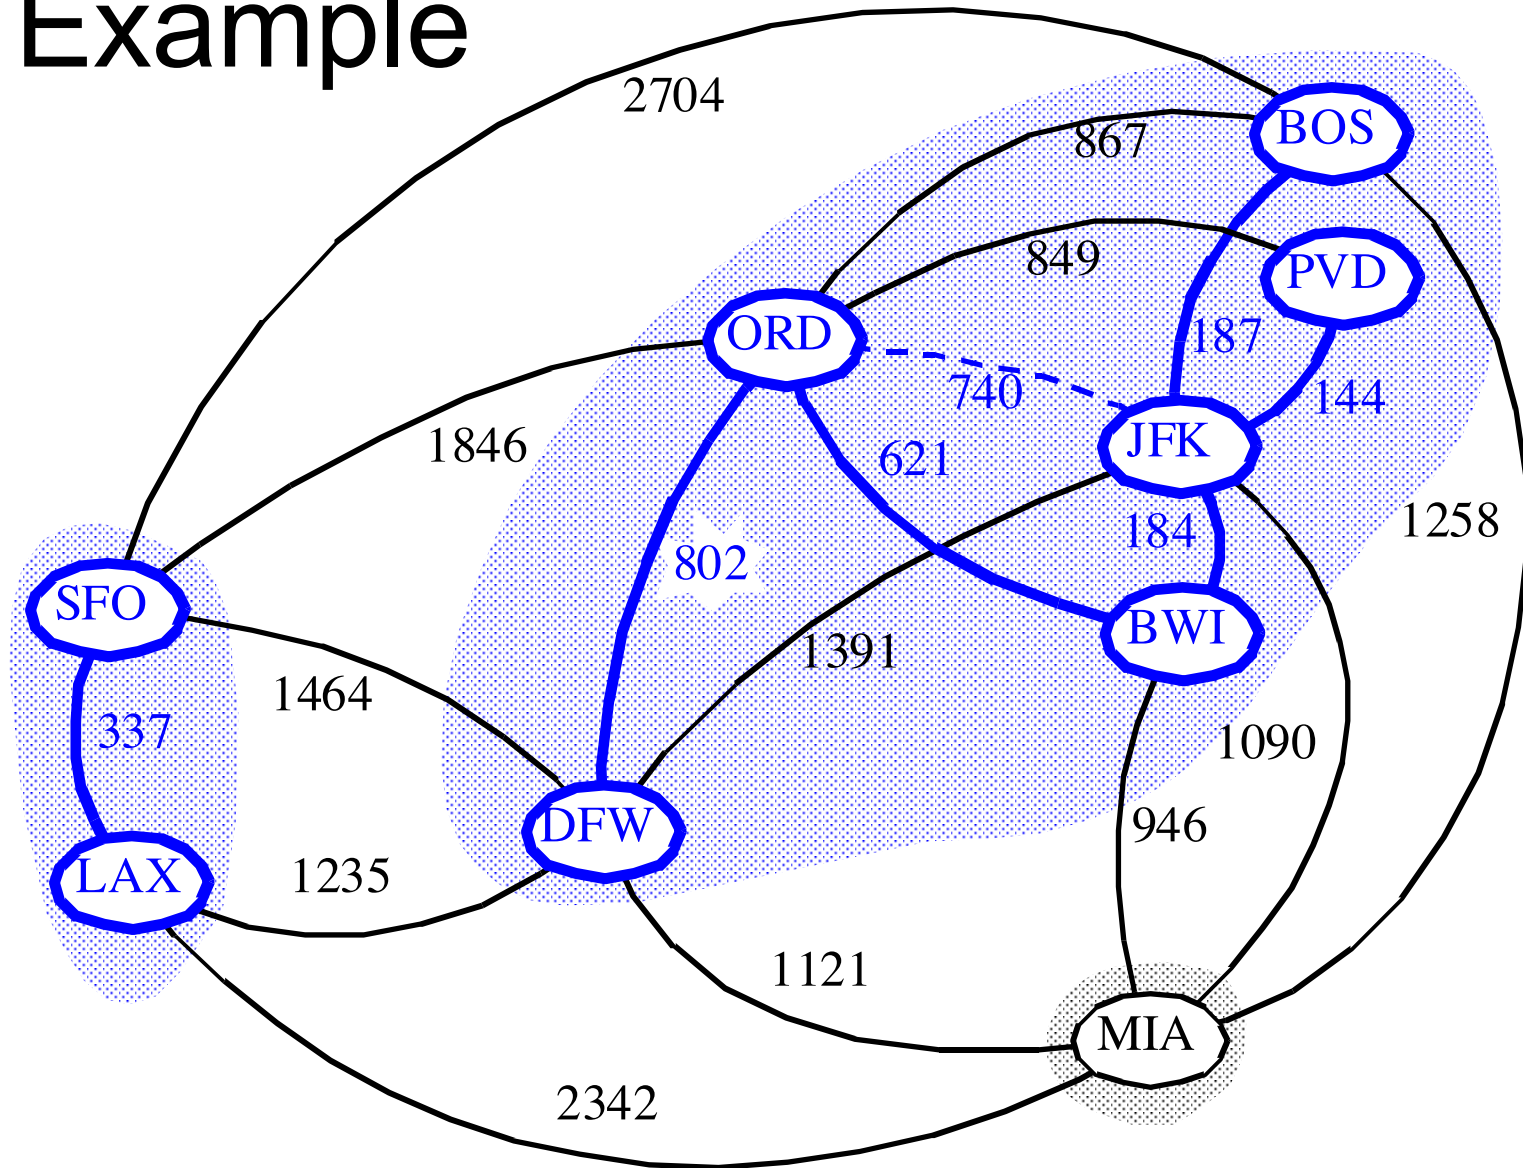
Other edges and weights include:

- SFO to ORD (2704)
- SFO to DFW (1846)
- SFO to LAX (1464)
- LAX to DFW (1235)
- DFW to ORD (802)
- DFW to MIA (1121)
- ORD to BOS (849)
- ORD to JFK (740)
- BOS to PVD (867)
- PVD to JFK (187)
- JFK to BWI (184)
- BWI to MIA (946)
- MIA to BOS (1258)

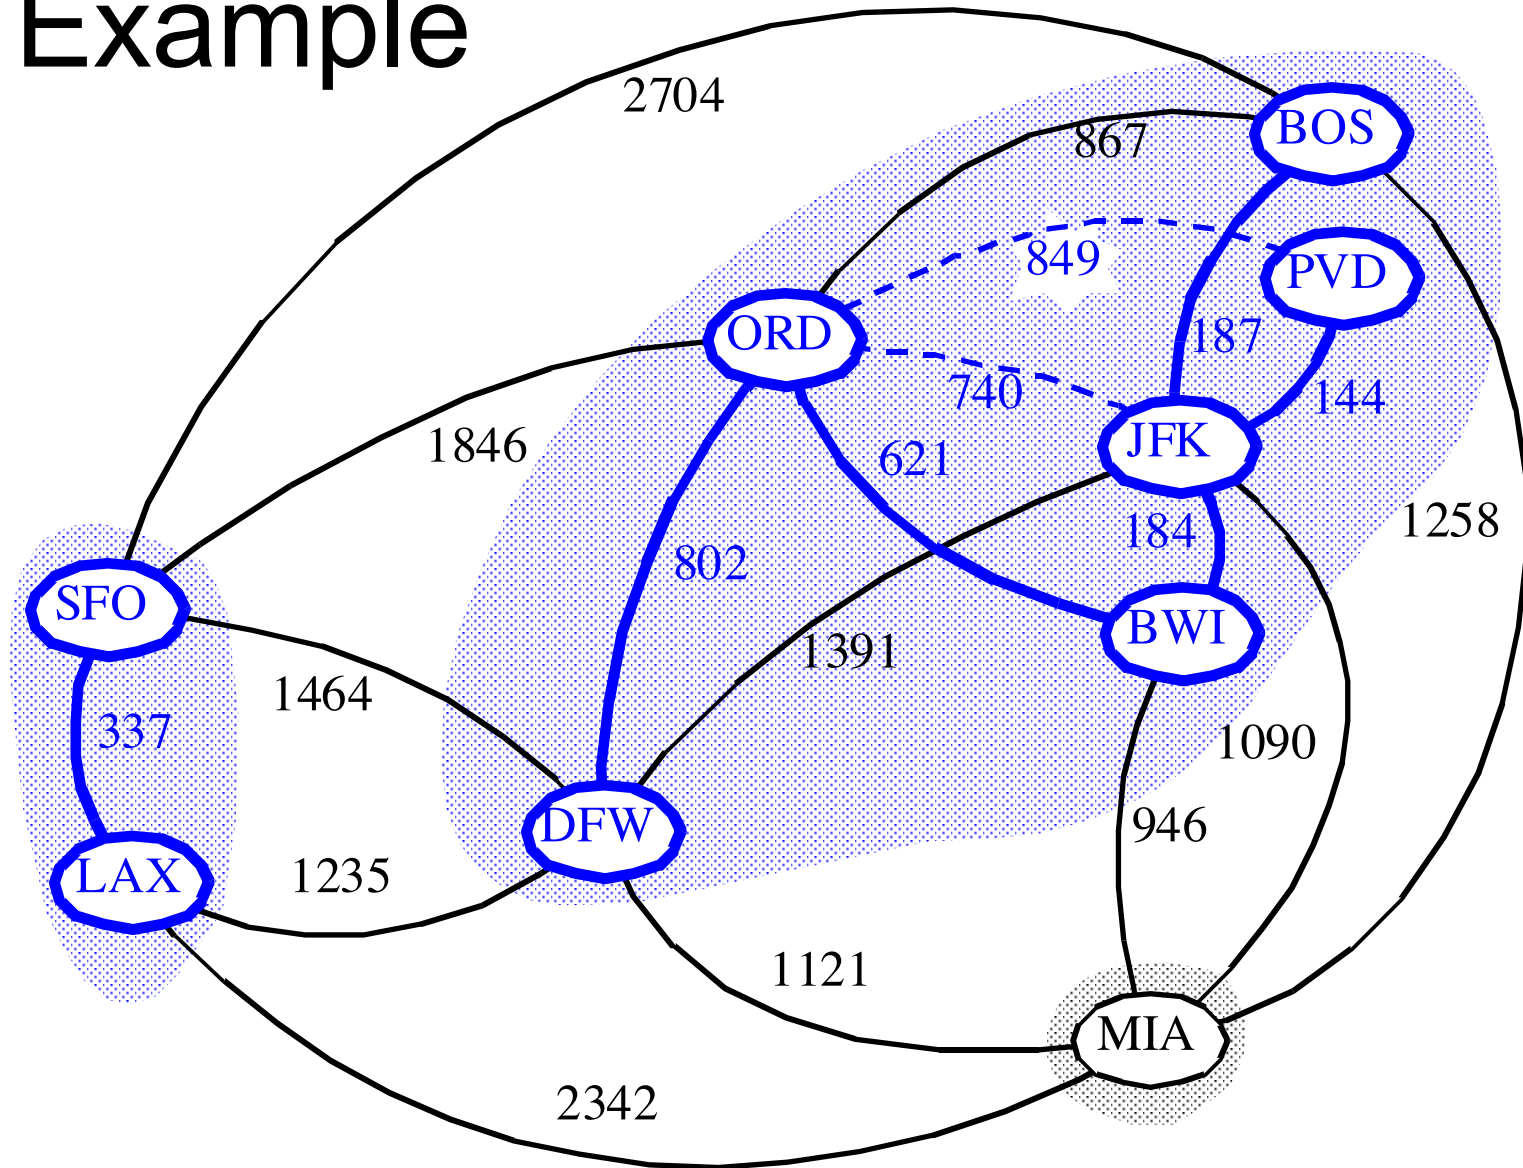
Example



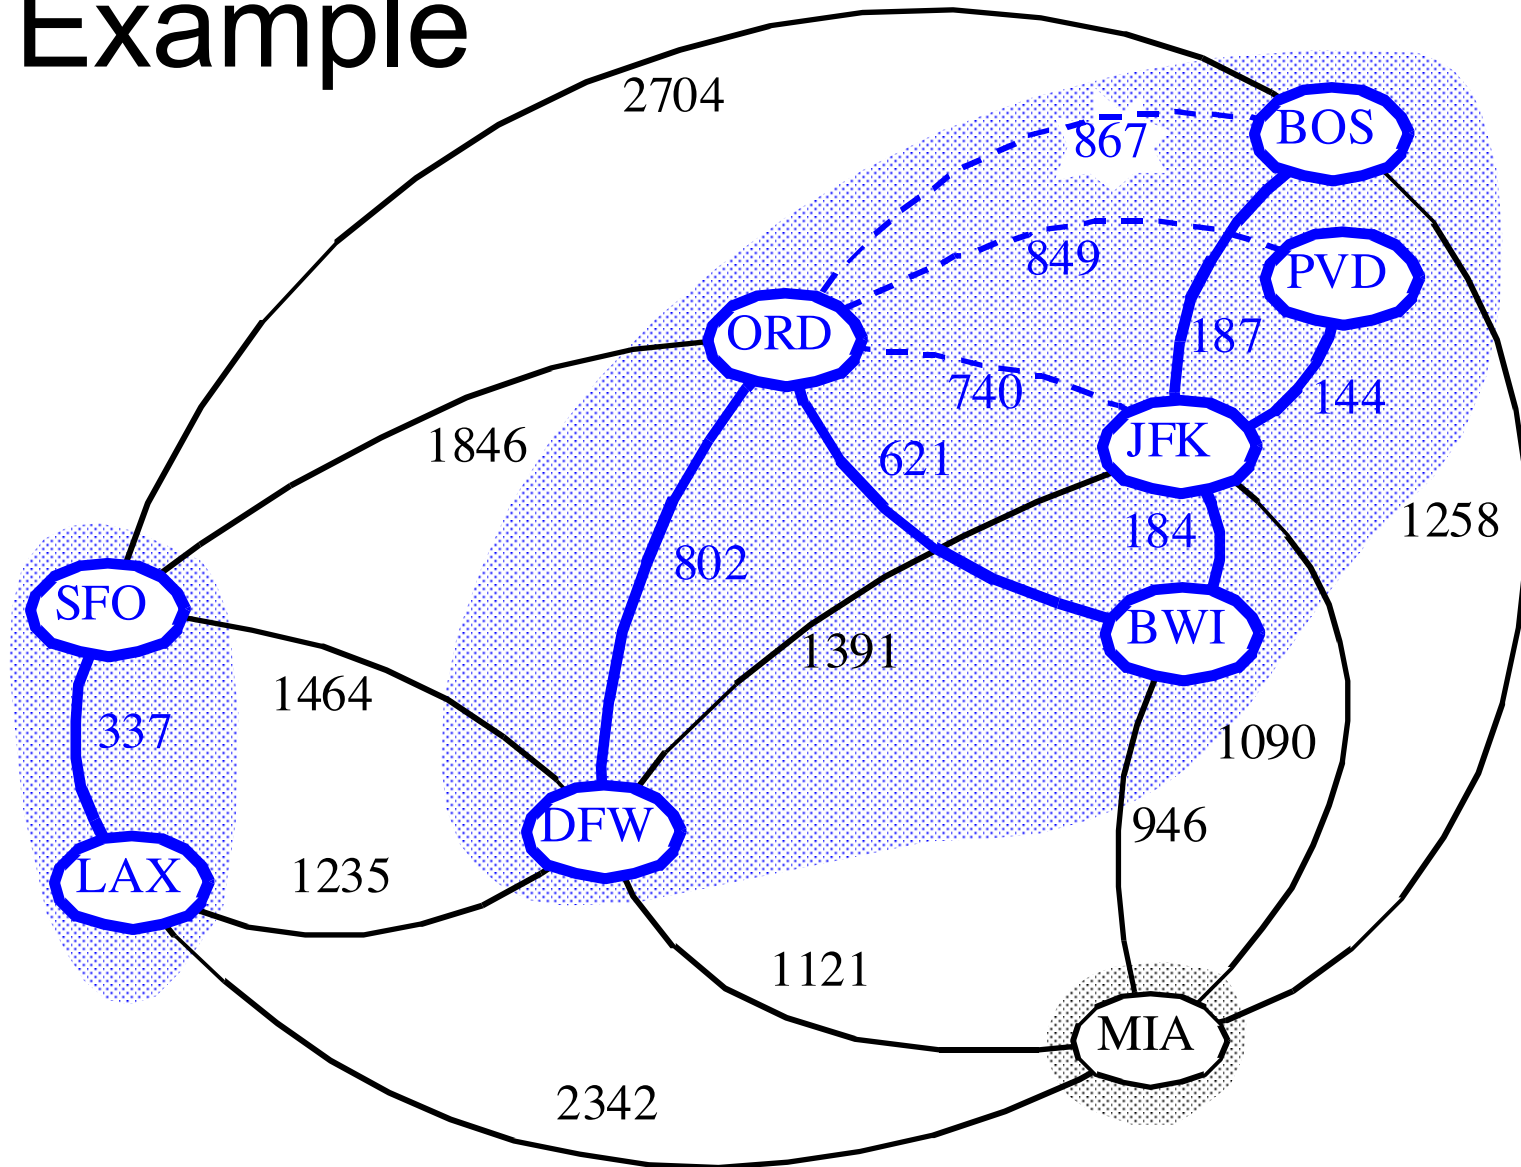
Example



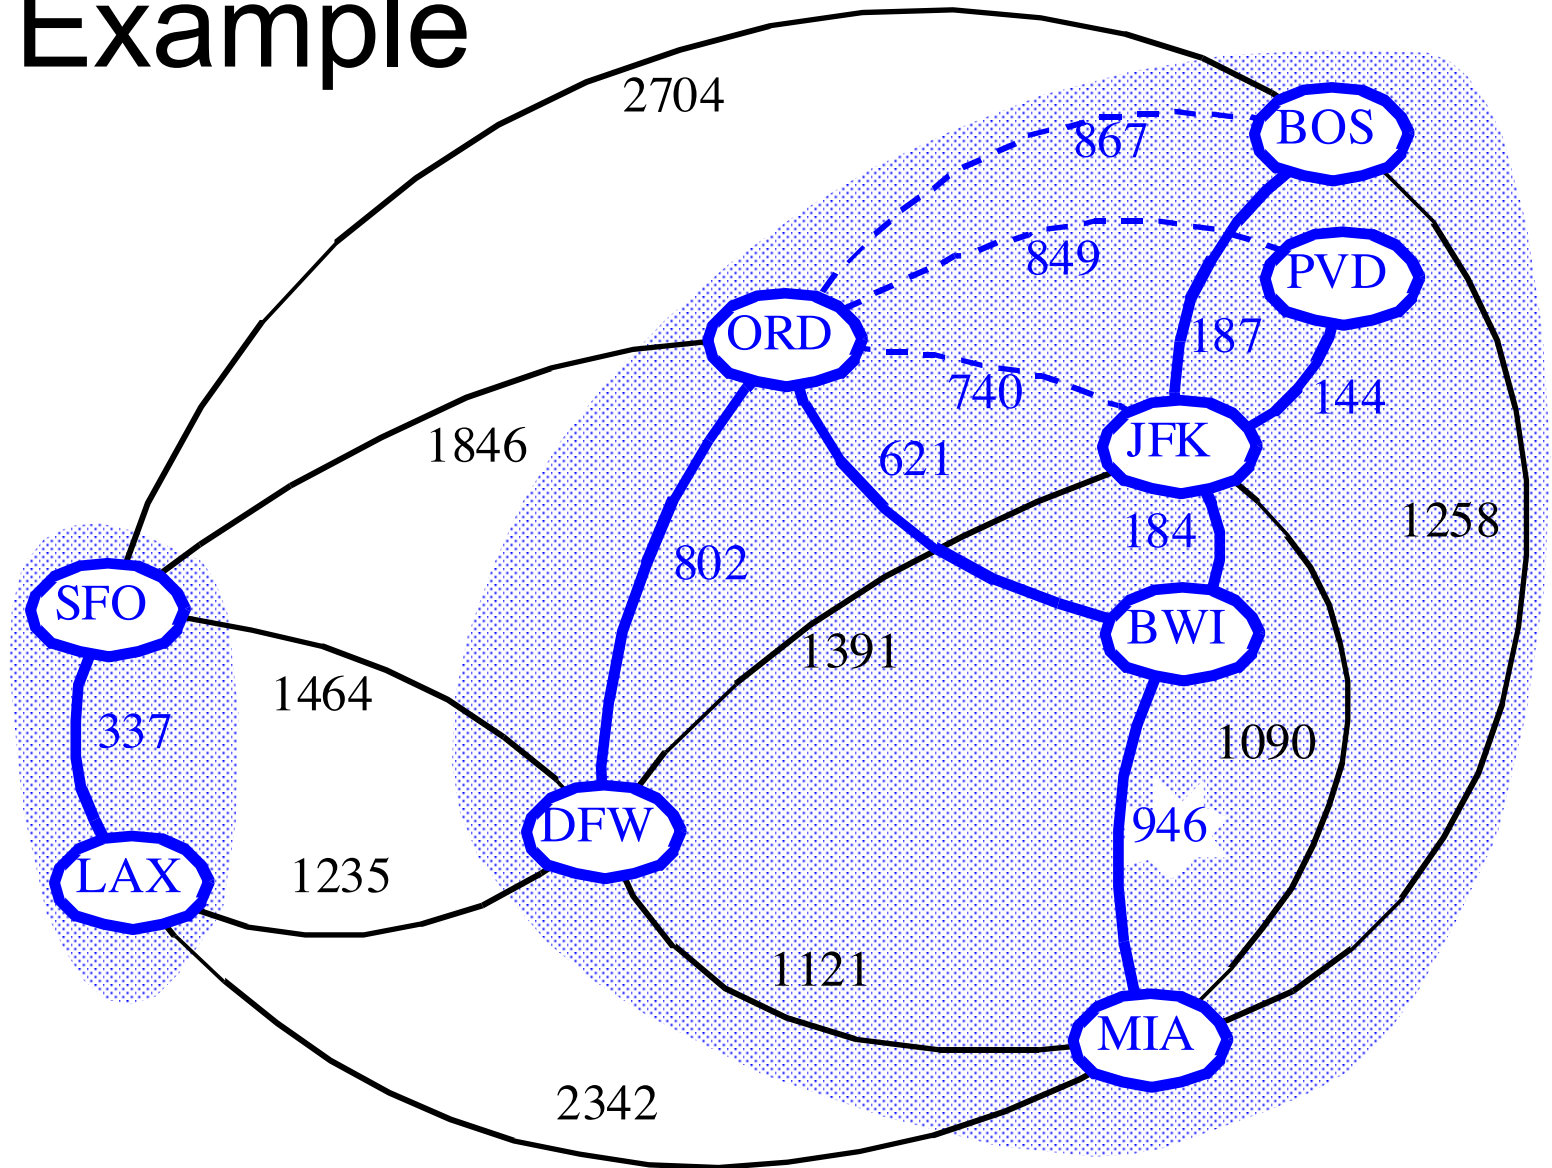
Example



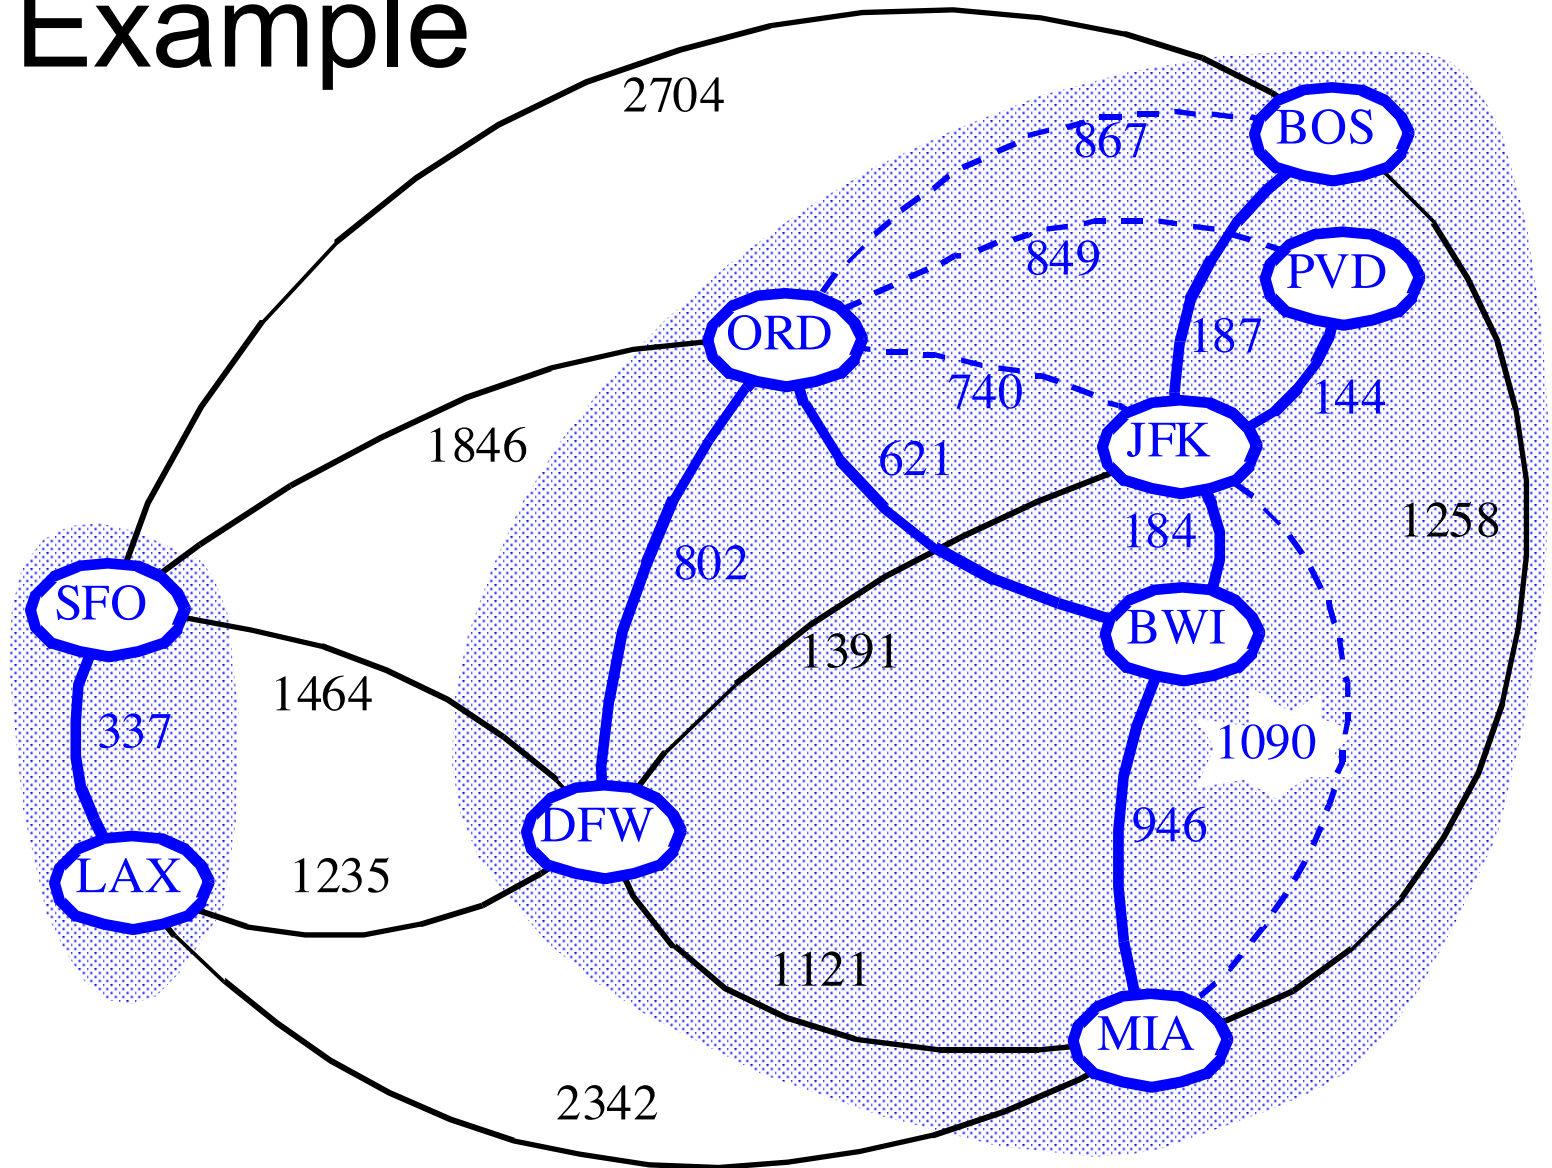
Example



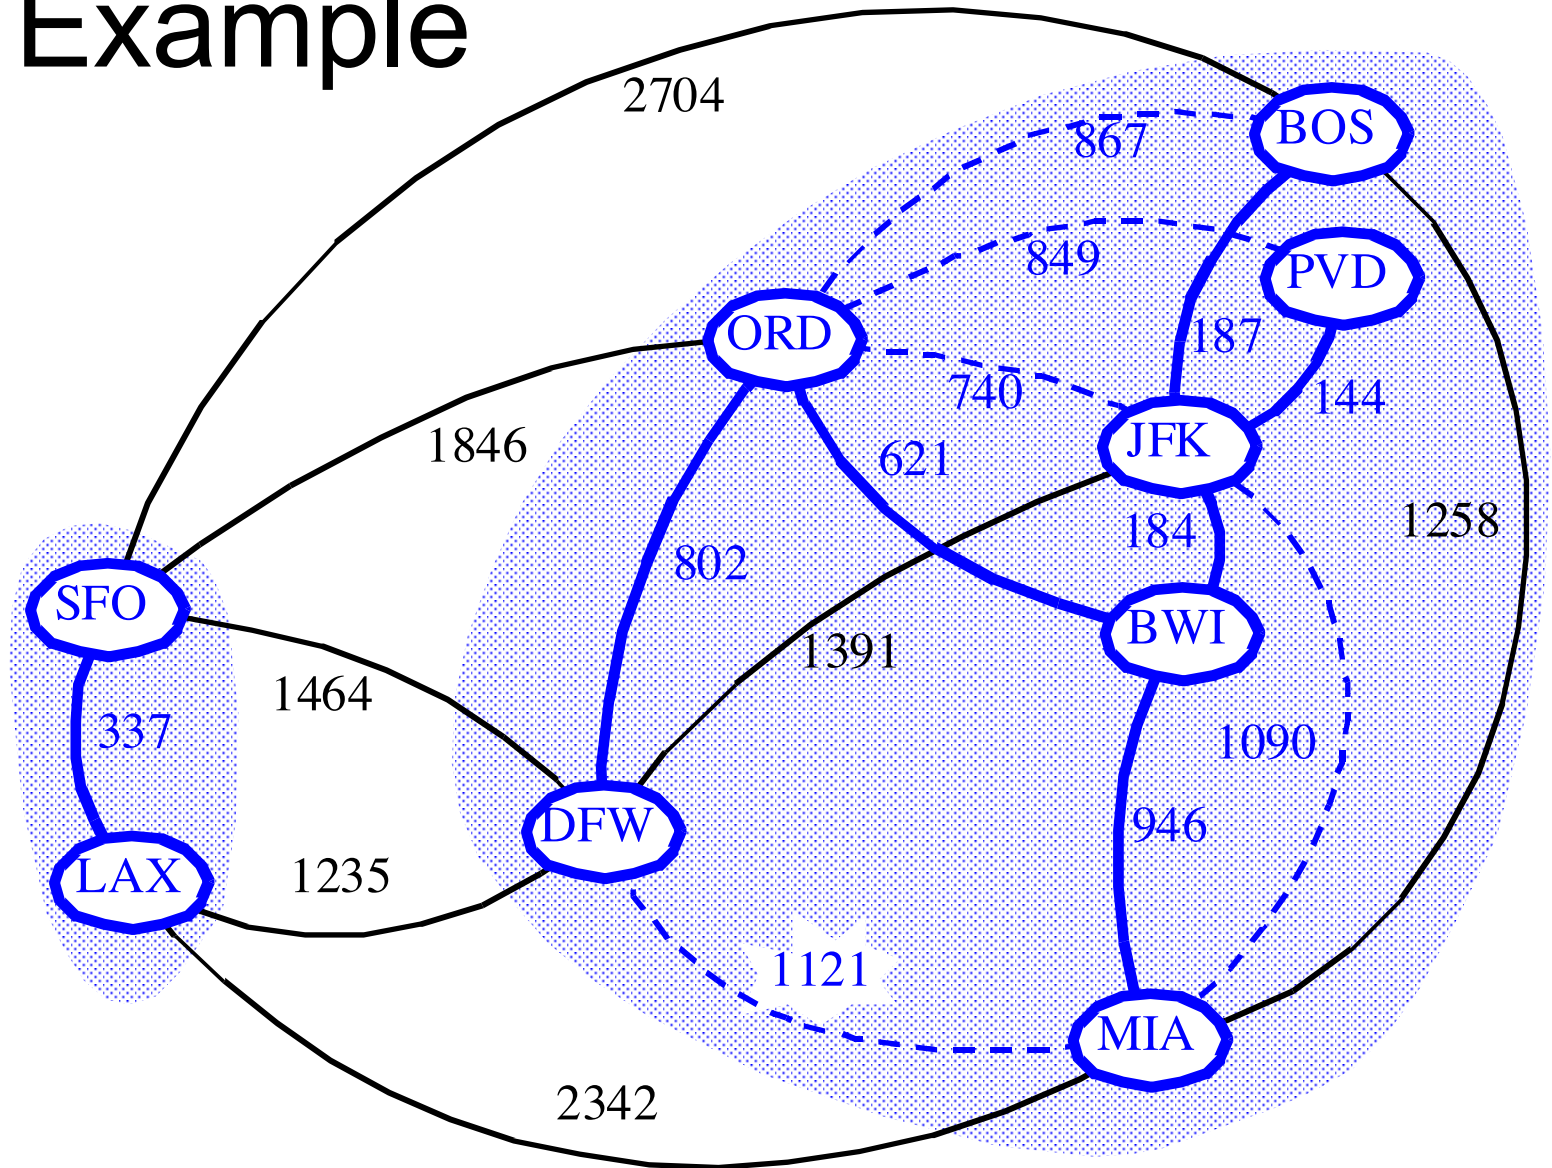
Example




Example



Example



Example



Minimum Spanning Tree

Baruvka's Algorithm

Like Kruskal's Algorithm, Baruvka's algorithm grows many “clouds” at once.


Algorithm *BaruvkaMST*(G)

```
 $T \leftarrow V$  {just the vertices of  $G$ }  
while  $T$  has fewer than  $n-1$  edges do  
  for each connected component  $C$  in  $T$  do  
    Let edge  $e$  be the smallest-weight edge from  $C$  to another component in  $T$ .  
    if  $e$  is not already in  $T$  then  
      Add edge  $e$  to  $T$   
return  $T$ 
```

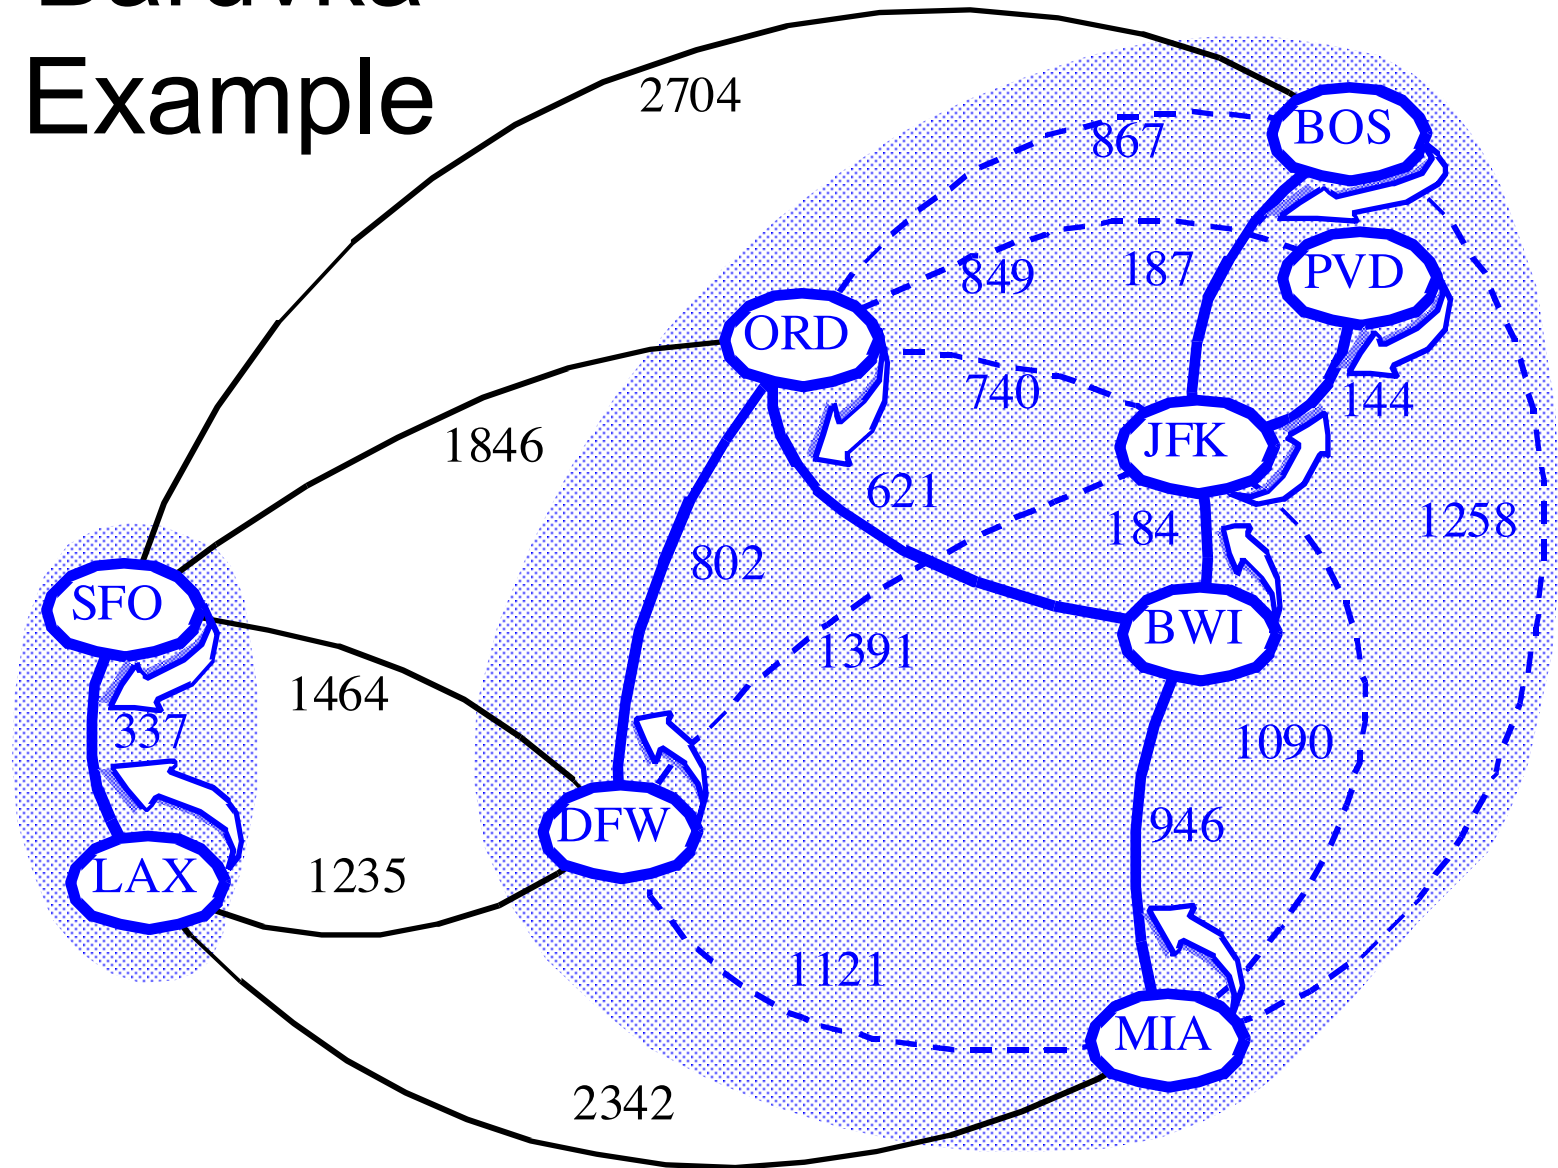
Each iteration of the while-loop halves the number of connected components in T .

- The running time is $O(m \log n)$.

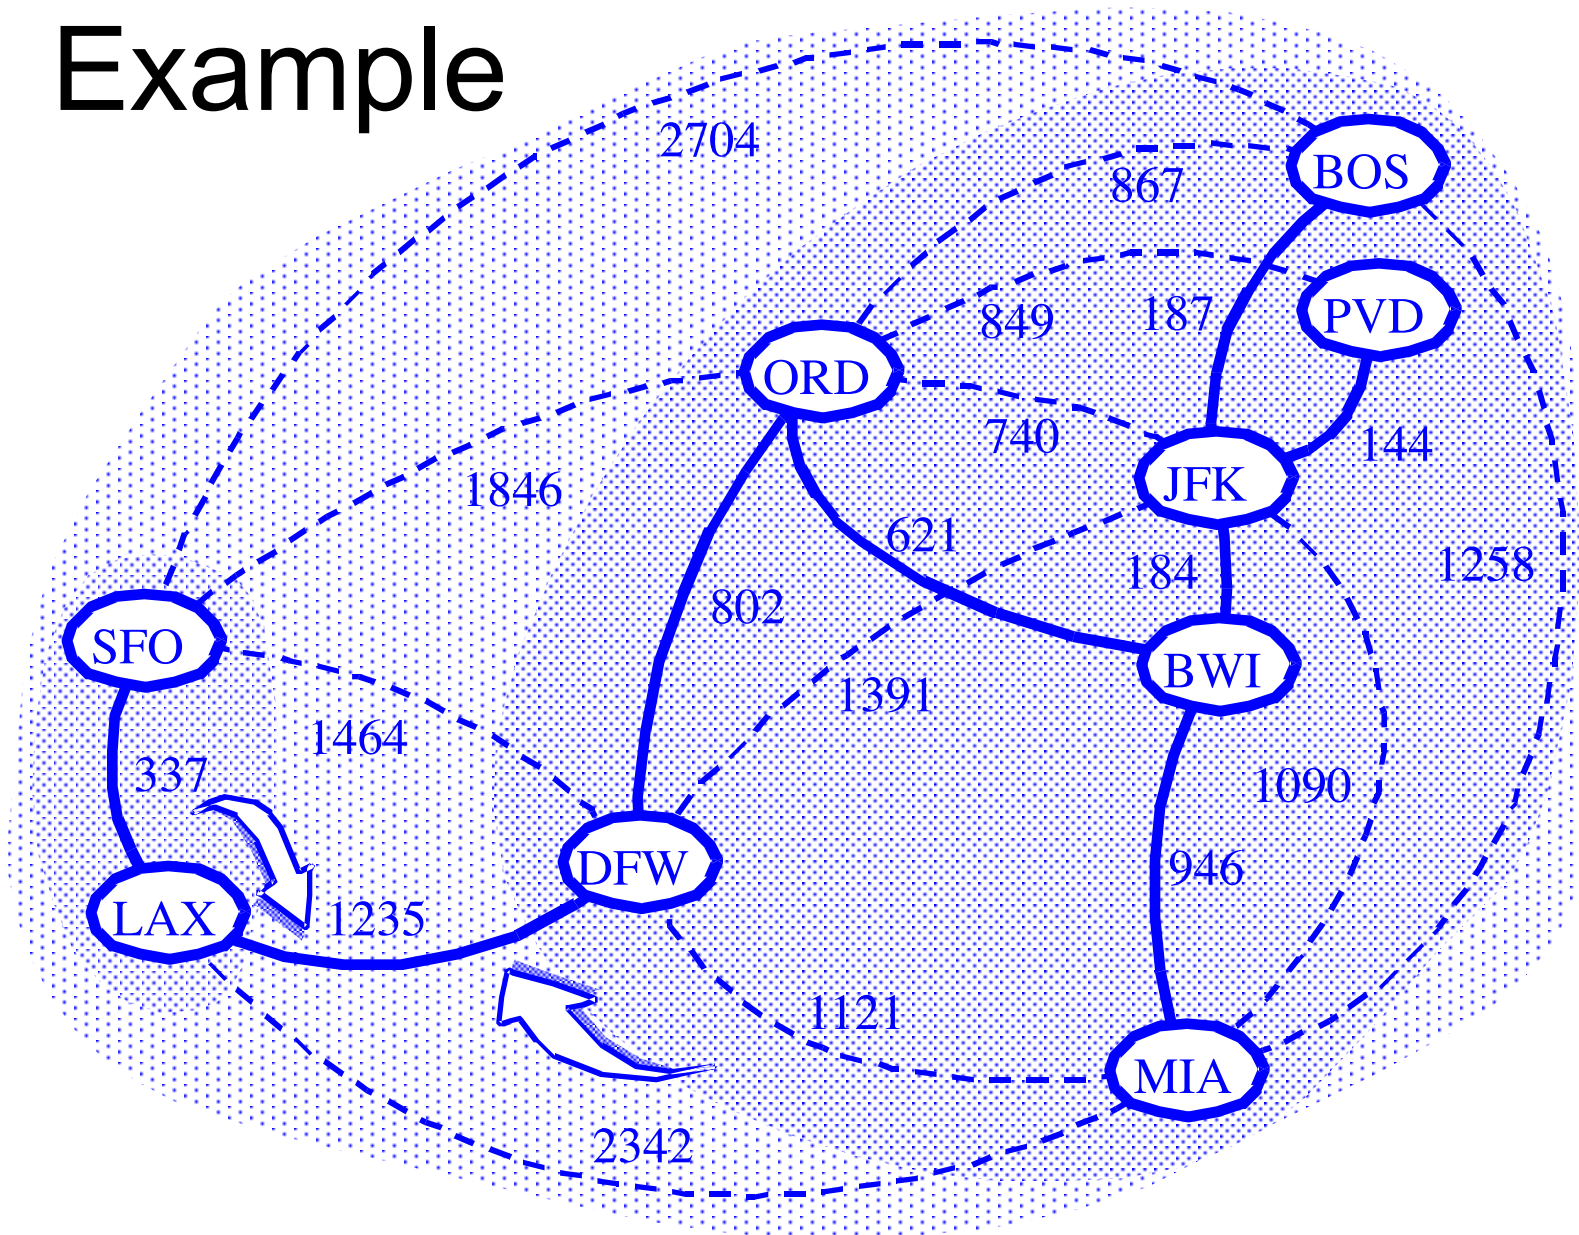
Baruvka Example



Baruvka Example



Example



Other – TSP

There's a long list of cities that Santa Claus needs to visit, and he only has from now until Christmas to figure out a good route to take.

Give an **efficient** algorithm that will come up with a guaranteed short route starting from the North Pole that will visit every city and come back to the North Pole again.

Note: You might not find the shortest route, but you can guarantee that you will always come *close* to the shortest route. This is known as an **approximation algorithm**. How closely does your algorithm approximate the optimal route?

